Computational Microsensitometry: A Physics-Based Pipeline for the Emulation of Kodak Vision3 250D Motion Picture Negative Film

1. Introduction: The Imperative for Physics-Based Emulation

The transition from photochemical imaging to digital acquisition has fundamentally reshaped the aesthetic and technical landscape of motion picture production. While digital sensors offer distinct advantages in terms of linearity, signal-to-noise ratio, and workflow efficiency, the specific non-linear radiometric response, complex chemical interactions, and stochastic structural characteristics of silver halide film remain a benchmark for "cinematic" image quality. Current industrial practices for film emulation predominantly rely on three-dimensional Color Lookup Tables (3D LUTs). While efficient for global color grading, LUT-based approaches are inherently limited: they map input code values to output code values based on static photometric measurements, ignoring the spatial and temporal dynamics that define the "film look." A LUT cannot simulate optical blooming, which depends on neighborhood luminance; it cannot replicate chemical adjacency effects (Mackie lines), which depend on local contrast gradients; nor can it synthesize signal-dependent granularity that interacts organically with image detail.

To bridge this gap, this report proposes a rigorous, physics-based image processing pipeline designed to emulate Kodak Vision3 250D (5207/7207) Color Negative Film. This approach, termed "Computational Microsensitometry," synthesizes classical photographic theory—derived from the seminal works of Mees, James, Todd, Zakia, Hunt, Friedman, and Ross—with contemporary differentiable optimization techniques. By modeling the distinct stages of the photographic process—spectral exposure, optical transport, latent image formation, chemical development, and granular reconstruction—we establish a framework that transcends empirical mimicry, offering a predictive simulation of the physical medium.

The pipeline is structured into seven distinct yet interconnected modules, each addressing a specific physical phenomenon identified in the literature:

1. **Spectral Exposure Calculation:** Modeling the interaction of scene irradiance with the spectral sensitivity of silver halide crystals.
2. **Optical Scattering:** Simulating the diffusion of light within the turbid emulsion and halation artifacts via Fast Fourier Transforms (FFT).
3. **Chemical Kinetics:** Modeling reaction-diffusion dynamics to reproduce adjacency

effects (edge enhancement).
4. **Sensitometric Tone Mapping:** applying the characteristic H&D curves specific to Vision3 250D.
5. **Colorimetric Transformation:** Simulating the subtractive mixing of Cyan, Magenta, and Yellow dyes and correcting for unwanted absorptions.
6. **Granularity Synthesis:** Employing Neural Film Grain Rendering (GrainNet) conditioned on optical density.
7. **Parameter Optimization:** Utilizing differentiable programming to fit simulation parameters to empirical datasheet constraints.

2. Module I: Spectral Exposure and Latent Image Physics

The fidelity of any film emulation pipeline is contingent upon the accuracy of its initial radiometric input. Unlike digital sensors, which typically employ a Bayer filter array with specific spectral transmission characteristics, photographic emulsions utilize layers of silver halide crystals sensitized to overlapping spectral bands. The emulation must therefore begin not with RGB pixel values, but with an estimation of the Spectral Power Distribution (SPD) of the incident light.

2.1 Silver Halide Sensitivity and Latent Image Formation

The fundamental mechanism of image capture in film is the formation of the latent image. As detailed in *The Theory of the Photographic Process* [1], this process is governed by the Gurney-Mott theory, where the absorption of photons by silver halide crystals generates photoelectrons. These electrons are trapped at sensitivity centers (often silver sulfide specks formed during chemical sensitization), where they attract interstitial silver ions to form metallic silver atoms.[1]

For Kodak Vision3 250D, a daylight-balanced film (5500K), the emulsion is constructed of multiple layers sensitized to different spectral regions. The "spectral sensitivity" ($S(\lambda)$) is defined as the reciprocal of the energy (erg/cm^2) required to produce a specified density (usually 1.0 above minimum density).[3]

Data Analysis of Vision3 250D Sensitivity:

- **Blue-Sensitive Layer:** This layer is inherently sensitive to blue and ultraviolet radiation. In Vision3 250D, it peaks near 440nm. A yellow filter layer is positioned below it to prevent blue light from exposing the lower layers, a critical structural component that digital sensors lack.[3]
- **Green-Sensitive Layer:** Orthochromatically sensitized using cyanine dyes, peaking around 550nm.
- **Red-Sensitive Layer:** Panchromatically sensitized, peaking near 640nm and extending

to approximately 680nm.

The pipeline calculates the "actinic exposure" H_c for each layer $c \in \{R, G, B\}$ by integrating the scene spectral radiance $L(\lambda)$ against the film's spectral sensitivity $S_c(\lambda)$:

$$H_c = \int_{\lambda_{min}}^{\lambda_{max}} L(\lambda) S_c(\lambda) d\lambda$$

This integration accounts for metamerism failures unique to film. For example, two colors that appear identical to a digital sensor (metamers) might expose the film differently due to the specific shape of the film's sensitivity curves, particularly in the deep red or cyan regions where digital filters often have sharper cutoffs.

2.2 Reciprocity Law Failure (RLF)

A critical deviation from ideal physics in silver halide emulsions is Reciprocity Law Failure. The Bunsen-Roscoe law states that the photochemical effect depends only on the total energy ($H = I \times t$). However, Mees & James [1] explain that at very low intensities (long exposures) or very high intensities (short electronic flashes), the efficiency of latent image formation drops. This is due to the instability of single silver atoms (thermal decay) at low intensities and the competition for electrons at high intensities.[1]

Although Kodak Vision3 250D is engineered for robust reciprocity (requiring no adjustment for exposures from 1/1000 sec to 1 sec [4]), accurately modeling this phenomenon is essential for emulating long-exposure aesthetics where shadow detail may be lost disproportionately. We model this using the Schwarzschild exponent p:

$$H_{eff} = I \cdot t^p$$

For Vision3 250D, p typically deviates from 1.0 only at exposures exceeding 1 second, necessitating a slight exposure increase and color compensation filter (e.g., CC10R) to correct for the differential failure rates of the RGB layers.[4]

3. Module II: Optical Scattering and Halation

Before the latent image is chemically fixed, the photons must traverse the physical medium of the emulsion. This medium is turbid, consisting of silver halide microcrystals suspended in gelatin. The propagation of light through this medium is subject to scattering, reflection, and absorption, processes that fundamentally define the image structure and "sharpness" of the film.

3.1 Turbid Medium Theory and Point Spread Functions

As described in Chapter 20 of *The Theory of the Photographic Process* [1], the optical properties of the emulsion can be modeled using radiative transfer theory. Light entering the emulsion is scattered by the silver halide grains. The intensity distribution of a point source imaged onto the film spreads out, a phenomenon quantified by the **Optical Point Spread Function (PSF)**.

For Vision3 250D, which incorporates "Sub-Micron Technology" to reduce grain and improve sharpness [5], the scattering is minimized but still present. Unlike the ideal diffraction-limited spot of a digital sensor, the film PSF typically exhibits a narrow central peak (forward scattering) surrounded by a wider skirt (multiple scattering). We approximate this optical spread using a Lorentzian distribution, which models the "long tails" of scattering in turbid media better than a Gaussian distribution [1]:

$$P_{scatter}(r) = \frac{1}{\pi \gamma (1 + (r/\gamma)^2)}$$

Here, γ is a scattering coefficient derived from the Modulation Transfer Function (MTF) curves provided in the film datasheet.[5] The MTF is the Fourier transform of the PSF; thus, by fitting the spectral response of our PSF model to the measured MTF data, we can determine the physical scattering parameters of the emulation.

3.2 Halation and Bloom Physics

Halation is a specific optical defect that has become a beloved aesthetic characteristic of analog film. It occurs when light penetrates the emulsion layers and the anti-halation backing (rem-jet), reflects off the interface between the film base and the air (or pressure plate), and re-enters the emulsion from the rear.6

Because the red-sensitive layer is typically the bottom-most layer (closest to the base) and because long-wavelength red light is scattered less and absorbed less by the upper layers, the reflected light predominantly exposes the red layer. This results in the characteristic reddish-orange halo around bright highlights.7

Bloom, in the context of film, refers to the local diffusion of light in the immediate vicinity of high-intensity regions, largely due to the turbidity of the emulsion and lens flare.[6]

3.3 Mathematical Modeling via FFT

To simulate these effects efficiently, we employ the Convolution Theorem. Calculating the convolution of a large halation kernel in the spatial domain is computationally expensive ($O(N^2 M^2)$). Utilizing the Fast Fourier Transform (FFT) reduces this to $O(N^2 \log N)$.

The total optical PSF, P_{total}, is modeled as a weighted sum of the scattering PSF and the halation PSF:

$$P_{total}(x, y, \lambda) = w_{scatter} P_{scatter}(x, y) + w_{halation}(\lambda) P_{halation}(x, y)$$

- **$P_{scatter}$:** Modeled as a narrow Lorentzian or Gaussian function representing turbidity.
- **$P_{halation}$:** Modeled as a broad annular (ring-shaped) function or wide Gaussian, offset by the thickness of the acetate base.
- **$w_{halation}(\lambda)$:** A wavelength-dependent weight. For Vision3 250D, this weight is maximal for the Red channel and minimal for Blue, simulating the preferential back-scattering of red light.[8]

The simulation proceeds by transforming the exposure image H and the PSF into the frequency domain:

$$H_{optical} = \mathcal{F}^{-1} \left(\mathcal{F}(H) \cdot \mathcal{F}(P_{total}) \right)$$

This operation introduces the characteristic "glow" and red-fringing around highlights before the image is processed sensitometrically, ensuring that the halation reacts non-linearly to the film's tone curve.

Feature	Physical Cause	Emulation Method
Sharpness Loss	Light scattering by AgX grains	Convolution with Lorentzian PSF
Halation	Reflection from base/air interface	Convolution with Red-weighted Annular PSF
Bloom	Emulsion turbidity & lens flare	Convolution with narrow Gaussian PSF

4. Module III: Chemical Adjacency Effects (Reaction-Diffusion)

Following exposure, the film undergoes chemical development (ECN-2 process for Vision3). This is not a simple pixel-wise amplification of the latent image; it is a complex reaction-diffusion system that gives rise to **adjacency effects**, most notably the Eberhard effect and Mackie lines.

4.1 Reaction-Diffusion Dynamics

As described by Friedman & Ross in Mathematical Models in Photographic Science 1 and Chapter 10 of their text 1, the development process involves the diffusion of developer into the gelatin and the counter-diffusion of reaction by-products.

The reduction of silver halide releases halide ions (bromide and iodide) which act as development inhibitors.

$$AgBr + Developer \rightarrow Ag + Oxidized Developer + Br^- (Inhibitor)$$

In areas of high exposure (highlights in the negative), a large concentration of inhibitor is produced. These inhibitors diffuse laterally into adjacent areas.

- **High Density Area:** High inhibitor concentration retards development.
- **Low Density Area:** Low inhibitor concentration allows full development.
- **The Border (Edge):** At the boundary between a high-density and low-density region, inhibitors from the high-density side diffuse into the low-density side, but fresh developer from the low-density side diffuses into the edge of the high-density area.

4.2 Eberhard and Mackie Lines

This dynamic creates two distinct artifacts:

1. **Mackie Lines:** A bright line (increased density) appears just inside the high-density edge (where fresh developer enters), and a dark line (decreased density) appears just outside the edge (where inhibitors accumulate).[1]
2. **Eberhard Effect:** Small high-density spots develop to a higher density than large areas of the same exposure because inhibitors can diffuse away from the small spot in all directions, whereas in a large area, they accumulate in the center.[9]

4.3 The Chemical Spread Function (CSF)

Nelson 1 introduces the Chemical Spread Function, $b(x)$, to quantify this lateral inhibition. Unlike the Optical PSF which spreads light, the Chemical Spread Function spreads inhibition. We model the effective density $D_{adj}(x, y)$ as a function of the nominal density $D(x, y)$ modulated by a blurred version of itself (representing the inhibitor cloud):

$$D_{adj}(x, y) = D(x, y) + \gamma_{chem} \left(D(x, y) - D(x, y) * b(x, y) \right)$$

where $*$ denotes convolution, $b(x, y)$ is the chemical spread function (typically exponential), and γ_{chem} controls the magnitude of the adjacency effect.

This operation is mathematically equivalent to an unsharp mask. In the frequency domain, this acts as a high-pass filter, boosting mid-to-high spatial frequencies. This chemical sharpening counteracts the optical blurring (low-pass filter) described in Module II. The interplay between optical blur and chemical sharpening is the "secret sauce" of film's acutance—it appears sharp (high edge contrast) yet smooth (no aliasing).

The parameters for the CSF are derived from the Modulation Transfer Function (MTF) curves of Vision3 250D.[5] While optical scattering degrades MTF at high frequencies, chemical adjacency boosts it. The measured MTF is the product of the Optical MTF and the Chemical MTF.

5. Module IV: Sensitometric Tone Mapping

The translation of developed silver mass into optical density is governed by the characteristic curve of the emulsion, known as the **H&D Curve** (Hurter and Driffield) or **D-Log E Curve**. This relationship is inherently non-linear and fundamental to the film's dynamic range compression.

5.1 The D-Log E Curve Structure

Todd & Zakia [1] analyze the characteristic curve in three distinct regions, which we must model parametrically:

1. **Toe:** The region of low exposure where the gradient gradually increases. This defines the shadow rolloff. Vision3 250D has a relatively long toe, allowing for smooth shadow detail compression without harsh clipping.
2. **Straight Line:** The linear region where density is proportional to log exposure. The slope of this line is **Gamma** (γ), which measures the contrast of the film. For color negative film like 5207, γ is typically around 0.50 to 0.60.[10]
3. **Shoulder:** The high-exposure region where the response saturates. Vision3 250D is renowned for its "extended highlight latitude".[11] Unlike digital sensors that clip hard at saturation, the Vision3 shoulder rolls off very gradually, retaining detail for 2+ stops beyond the linear range.[11]

5.2 Parametric Modeling and Optimization

To emulate this, we do not use a simple lookup table. Instead, we fit a continuous, differentiable function to the sensitometric data points provided in the Kodak datasheet.5
A suitable parametric model is a Generalized Logistic Function or a Cubic Spline. The function $f_{H\&D}$ maps the log-actinic exposure $h = \log_{10}(H)$ to density D:

$$D = D_{min} + \frac{D_{max} - D_{min}}{(1 + e^{-k(h - h_0)})^\nu}$$

where parameters k, h_0, ν control the slope (gamma) and the asymmetry of the toe and shoulder.
Using differentiable optimization (discussed in Module VII), we adjust these parameters so that the simulated curve minimizes the Mean Squared Error (MSE) against the datasheet curves for the Red, Green, and Blue layers.
Important Note on Gamma: The "Gamma" of the negative film is low (~0.55). However, the final viewed image requires a system gamma of roughly 1.0 to 1.6 depending on viewing conditions (dark vs. dim surround).[1] The pipeline must therefore include a "print" emulation step (or Scan-to-Scene transform) to expand the contrast of the negative to a viewable range, mimicking the response of print film (e.g., Kodak Vision Color Print Film 2383).

6. Module V: Colorimetric Transformation

Once the silver densities are established, they must be converted into dye densities. In color negative film, the metallic silver is bleached out, leaving behind cyan, magenta, and yellow dyes formed by couplers.

6.1 Subtractive Color Mixing and Unwanted Absorptions

As Hunt explains in *The Reproduction of Colour* [1], ideal subtractive dyes would absorb only their complementary bandwidths (e.g., Cyan absorbs Red). Real dyes, however, have "unwanted absorptions."

- Cyan dye absorbs some Green and Blue light.
- Magenta dye absorbs some Blue light.
- Yellow dye is relatively pure but still has minor imperfections.

This cross-coupling means that density in one channel affects the transmission in others. The relationship between analytical densities (amount of dye, C, M, Y) and integral densities (measured absorption D_R, D_G, D_B) is linear, governed by Beer's Law [1]:

$$\begin{bmatrix} D_R \\ D_G \\ D_B \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} \begin{bmatrix} C \\ M \\ Y \end{bmatrix} + \begin{bmatrix} D_{base,R} \\ D_{base,G} \\ D_{base,B} \end{bmatrix}$$

The off-diagonal elements of the matrix (a_{12}, a_{13}, etc.) represent the unwanted absorptions.

6.2 The Orange Mask (Colored Couplers)

To correct for these unwanted absorptions, color negative film utilizes colored couplers. The unreacted couplers remain in the film and have a color (yellow or pinkish) that complements the unwanted absorption of the dye they form. This results in the characteristic orange cast of the negative.1

In our pipeline, we simulate this by using the "Spectral Dye Density Curves" from the Vision3 250D datasheet 3 to populate the mixing matrix. The base density vector $$ accounts for the orange mask.

Since the pipeline output is typically a digital positive, we perform this transformation to generate the "Virtual Negative," and then apply a "Virtual Print" transform (inverting the matrix or using a print film LUT) to visualize the result. This ensures that the color crosstalk characteristic of film—which contributes to its specific palette—is mathematically preserved.

7. Module VI: Granularity and Texture Synthesis

Film grain is not additive Gaussian noise; it is a signal-dependent texture resulting from the stochastic distribution of silver halide crystals and the subsequent dye clouds.

7.1 Selwyn Granularity

Mees & James 1 define Selwyn Granularity (G) as the invariant product of the standard deviation of density fluctuations (σ_D) and the square root of the scanning aperture area (a):

$$G = \sigma_D \sqrt{2a}$$

For Vision3 250D, granularity is highly signal-dependent. In a negative, shadow areas (low density) on the print correspond to low density on the negative. However, grain perception is complex. The datasheet 5 indicates that granularity correlates with density.
Crucially, Vision3 250D uses "Dye Layering Technology" to reduce grain in underexposed areas (shadows).12 This means the relationship between signal intensity and noise variance is non-linear and specific to this stock.

7.2 Neural Film Grain Rendering (GrainNet)

Traditional procedural noise (Perlin or Simplex) fails to capture the morphological characteristics of dye clouds, which are clumpy and organic. We employ **GrainNet**, a neural network architecture proposed by Lesné et al..[13]

- **Architecture:** GrainNet utilizes a lightweight U-Net or ResNet-based convolutional neural network.
- **Conditioning:** The network takes two inputs: the "clean" grain-free image (from Module V) and a noise map. Critically, the network is **conditioned on the local density** of the image. This allows the network to modulate the grain structure (size and intensity) based on the H&D curve position, correctly simulating the signal-dependent nature of Vision3 grain.[13]
- **Texture-Aware Loss:** Training involves a texture loss (Gram matrix or spectral loss) to ensure the Power Spectral Density (PSD) of the synthesized grain matches real scans of Vision3 250D.

This module ensures that the "grain" is not just an overlay, but a structural component of the image that reacts to exposure, mirroring the physical reality of the emulsion.

8. Module VII: Differentiable Optimization and Parameter Fitting

The pipeline described above relies on numerous physical parameters: scattering coefficients γ, chemical diffusion lengths, H&D curve shape parameters, and dye matrix coefficients. Determining these manually is error-prone. We propose a **differentiable optimization framework** to solve this inverse problem.

8.1 The Optimization Problem

We define the entire emulation pipeline as a differentiable function $P(\theta, I)$, where

θ represents the set of physical parameters and I is the input image.
We aim to find the optimal parameters θ^* that minimize the difference between the pipeline's output and reference data (ground truth).

8.2 Implementation via Differentiable Programming

Using frameworks like JAX or PyTorch (or specialized tools like Tesseract 15), every step of the pipeline—from spectral integration to FFT convolution to curve application—is implemented as a differentiable operation.16
This allows us to compute the gradient of a Loss Function $\mathcal{L}$ with respect to the parameters θ:

$$\theta_{new} = \theta_{old} - \eta \nabla_\theta \mathcal{L}$$

8.3 Ground Truth and Loss Functions

We utilize the Kodak Vision3 250D datasheet and scanned reference charts as ground truth.

- **Sensitometric Loss:** We define a loss $\mathcal{L}_{curve}$ as the $L2$ distance between the pipeline's generated densities and the published characteristic curves.[5]
- **MTF Loss:** We compare the frequency response of the simulation (after Optical and Chemical modules) to the published MTF curves [5] to optimize the scattering γ and chemical boost parameters.
- **Perceptual Loss:** For grain and color, we use a perceptual metric (like LPIPS) or a texture loss to match the visual statistics of the simulation to scanned Macbeth charts exposed on 5207 stock.

This "End-to-End Camera Design" approach [18] ensures that the parameters are not merely theoretically sound but empirically tuned to the exact specification of the target film stock.

9. Conclusion

The pipeline presented herein represents a paradigm shift in film emulation, moving from empirical color mapping to **Computational Microsensitometry**. By systematically modeling the physics of the photographic process—spectral photon integration, optical scattering in turbid media, reaction-diffusion chemical kinetics, sensitometric non-linearities, and stochastic granular synthesis—we achieve a level of fidelity impossible with LUTs alone. The integration of differentiable optimization serves as the final bridge between theory and reality, allowing the mathematical models to be tightly calibrated against the empirical data of Kodak Vision3 250D. This framework not only preserves the aesthetic legacy of analog film but provides a robust, resolution-independent foundation for future archival and creative applications.

Table 1: Comparative Analysis of Emulation Modules

Module	Physical Phenomenon	Mathematical Model	Parameter Source (Vision3 250D)
I. Spectral	Photon Absorption / Quantum Yield	$H = \int L(\lambda)S(\lambda)d\lambda$	Spectral Sensitivity Curves [3]
II. Optical	Scattering & Halation	$H' = \mathcal{F}^{-1}(\mathcal{F}(H) \cdot \mathcal{F}(PSF))$	MTF Data (Low Freq) [5]
III. Chemical	Adjacency / Inhibitor Diffusion	$D_{adj} = D + \gamma(D - D*b)$	MTF Data (High Freq) / Edge Traces
IV. Sensitometric	D-Log E Response	Generalized Logistic Function	Characteristic Curves [5]
V. Colorimetric	Dye Absorption / Cross-talk	Matrix Multiplication (Beer's Law)	Spectral Dye Density Curves [3]
VI. Granularity	Stochastic Dye Cloud Distribution	Neural Texture Synthesis (GrainNet)	RMS Granularity Data [5]
VII. Optimization	Parameter Tuning	$\min_\theta	

Works cited

1. The Theory of the Photographic Process.pdf
2. Methodology and Curve Interpretation | Photrio.com Photography Forums, accessed December 18, 2025, https://www.photrio.com/forum/threads/methodology-and-curve-interpretation.189935/
3. KODAK VISION3 250D Color Negative Film 5207 / 7207 / SO-207 - MOnAMi, accessed December 18, 2025, https://monami.hs-mittweida.de/files/337/Anhang6_KODAK_VISION3_250D_Color.

pdf
4. EXPANDING THE VISION3 FILM PLATFORM—AND THE RANGE OF POSSIBILITIES IN CAPTURE AND POST. - Argentix.Ca, accessed December 18, 2025, https://www.argentix.ca/specs/V3-5207-en.pdf
5. KODAK VISION3 250D Color Negative Film 5207 / 7207, accessed December 18, 2025, https://www.kodak.com/content/products-brochures/Film/VISION3-250D-Technical-Data-EN.pdf
6. Bloom & Halation — Darkroom, accessed December 18, 2025, https://darkroom.co/blog/2025-12-08-bloom-halation
7. Halation is a visual effect that appear when shooting on a film as a red-orange halo near the contrasting boundaries of over-exposed areas, as well as a red glare in the mid-tones. Usually, Halation is produced around bright light sources. - Dehancer, accessed December 18, 2025, https://www.dehancer.com/learn/articles/halation-in-dehancer
8. Film Halation | Color.io User Guide, accessed December 18, 2025, https://www.color.io/user-guide/film-halation
9. Point spread function - Wikipedia, accessed December 18, 2025, https://en.wikipedia.org/wiki/Point_spread_function
10. What Does the Characteristic Curve or Density Curve Indicate - Baker Hughes, accessed December 18, 2025, https://www.bakerhughes.com/waygate-technologies/blog/what-does-characteristic-curve-or-density-curve-indicate
11. KODAK VISION3 250D - Color Negative Film 5207/7207, accessed December 18, 2025, https://www.kodak.com/content/pdfs/KODAK-VISION3-250D-5207-7207-brochure.pdf
12. VISION3 250D Color Negative Film 5207/7207 - Kodak, accessed December 18, 2025, https://www.kodak.com/en/motion/product/camera-films/250d-5207-7207/
13. Neural Film Grain Rendering - Eurographics Association, accessed December 18, 2025, https://diglib.eg.org/items/a0985d06-ee1c-402c-8d9f-827c7258d9bc
14. Our film grain rendering neural network architecture. δ corresponds to... | Download Scientific Diagram - ResearchGate, accessed December 18, 2025, https://www.researchgate.net/figure/Our-film-grain-rendering-neural-network-architecture-d-corresponds-to-the-grain-size-in_fig2_390921802
15. Pipeline-level differentiable programming for the real world - SciPy Proceedings, accessed December 18, 2025, https://proceedings.scipy.org/articles/kvfm5762
16. An overview of Differentiable Rendering | by Rémi B | Qarnot - Medium, accessed December 18, 2025, https://medium.com/qarnot/an-overview-of-differentiable-rendering-20ceb8d20cbe
17. Differentiable Compound Optics and Processing Pipeline Optimization for End-to-end Camera Design | Request PDF - ResearchGate, accessed December 18, 2025, https://www.researchgate.net/publication/352603602_Differentiable_Compound_

Optics_and_Processing_Pipeline_Optimization_for_End-to-end_Camera_Design
18. Differentiable Compound Optics and Processing Pipeline Optimization for End-to-end Camera Design - Princeton Computational Imaging Lab, accessed December 18, 2025, https://light.cs.princeton.edu/wp-content/uploads/2021/02/DeepCompoundOptics.pdf